

The Gems of Kubernetes' Latest Features

David Pech
2025



David Pech



Why bother learning about that hyped K-word?

"So you can have the same features as before on VMs."

... but of course this time *smarter*.

Online Pod Resizing

In-place vertical scaling without restarts

- Status: **Beta** in Kubernetes 1.34
- KEP: [1287](#)

The Problem Before

```
apiVersion: v1
kind: Pod
metadata:
  name: memory-hungry-app
spec:
  containers:
  - name: app
    image: my-app:v1
  resources:
    limits:
      memory: "2Gi" # <-- restart
      cpu: "1000m" # <-- restart
```

```
apiVersion: v1
kind: Pod
metadata:
  name: resizable-app
spec:
  containers:
  - name: app
    image: my-app:v1
    resources:
      limits:
        memory: "4Gi" # CHANGED
        cpu: "2000m" # CHANGED
  resizePolicy:
  - { resourceName: memory, restartPolicy: NotRequired }
  - { resourceName: cpu, restartPolicy: NotRequired }
```

Why Online Pod Resizing Matters

Use Cases:

- **Databases:** Scale memory up for heavy query periods
- **Auto-scaling:** Faster response to load changes

But:

- You probably don't care about CPU .limits
- Memory can scale only up
- Java

Pod-level Resource Allocation

Simplify with per-pod resource constraints

- Status: **Beta** since Kubernetes 1.34
- KEP: [2837](#)

```
apiVersion: v1
kind: Pod
metadata:
  name: database-pod
spec:
  containers:
    - name: postgres
      image: postgres:18
      resources:
        requests:
          cpu: "2000m"
          memory: "4Gi"
        limits:
          cpu: "2000m"
          memory: "4Gi"
    - name: exporter
      image: postgres_exporter
      resources: # <-- ??? WHAT TO PUT HERE?
        requests:
          cpu: "20m"
          memory: "4Mi"
        limits:
          cpu: "200m"
          memory: "40Mi"
```

```
apiVersion: v1
kind: Pod
metadata:
  name: database-pod
spec:
  resources: # Pod-level definition, shared between containers
    requests:
      cpu: "2000m"
      memory: "4Gi"
    limits:
      cpu: "2000m"
      memory: "4Gi"
  containers:
    - name: postgres
      image: postgres:18
    - name: exporter
      image: postgres_exporter
```

Why Pod-level Allocation Matters

Use Cases:

- Sidecars
- (Uneven containers)

Benefits: Easier to understand, Less waste, QoS class

But: Some tooling not ready yet, No in-place resize

VolumeSnapshots & VolumeGroupSnapshots

Simplified backup and restore across multiple volumes

- Status: VolumeSnapshots **Stable**, VolumeGroupSnapshot **Beta** in 1.34
- KEP: [3476](#)

The Problem Before

```
# Manual, error-prone backup process
kubectl exec mysql-pod -- mysqldump > backup.sql
kubectl cp mysql-pod:backup.sql ./local-backup.sql
```

```
# No coordination between related volumes
# Risk of inconsistent state across multiple volumes
# Manual restore process
```

And Now

```
# Snapshot for single PVC
apiVersion: snapshot.storage.k8s.io/v1
kind: VolumeSnapshot
metadata:
  name: mysql-snapshot-20241001
spec:
  volumeSnapshotClassName: csi-snapshotter
  source:
    persistentVolumeClaimName: mysql-pvc
# --> real data managed by VolumeSnapshotContent
```

And Now

```
# Multiple PVCs at the same time (multiple "mounts")
apiVersion: groupsnapshot.storage.k8s.io/v1beta1
kind: VolumeGroupSnapshot
metadata:
  name: app-group-snapshot
spec:
  volumeGroupSnapshotClassName: csi-group-snapshotter
  source:
    selector:
      matchLabels:
        app: database-cluster
# --> generates VolumeGroupSnapshotContent
```

Why VolumeSnapshots Matter

Database size	PGDATA volume size	WAL volume size	Snapshot full backup time	Object store full backup time	Snapshot recovery time	Object store recovery time
4.5 GB	8 GB	1 GB	1m 50s	9m 15s	31s	3m 29s
44 GB	80 GB	10 GB	20m 38s	1h 6m	27s	31m 59s
438 GB	800 GB	100 GB	2h 42m	9h 53m	48s	59m 51s
4381 GB	8000 GB	200 GB	3h 54m 6s	95h 12m 20s	2m 2s	10h 6m 17s

* Benchmarked using AWS EBS gp3 disks
* The test considers base backup recovery only, without WAL file recovery

Source (Kubecon NA, Chicago)

Why VolumeSnapshots Matter

Use Cases:

- **Database Backups:** Consistent point-in-time snapshots
- **Disaster Recovery:** Quick restoration from snapshots
- **Development/Testing:** Clone production data safely
- **Multi-volume Applications:** Coordinated snapshots across volumes

OCI ImageVolumes

Reduce image duplication and network overhead

- Status: **Beta** in Kubernetes 1.34
- KEP: [4639](#)

The Problem Before

```
FROM ai/gemma3:latest as model  
  
FROM myapp:latest AS app  
  
# Copy model files  
ADD --from=model model-file /mnt/model-file
```

```
docker build -t myapp:latest-with-model
```

The Problem Before

```
apiVersion: v1
kind: Pod
metadata:
  name: my-fancy-app
spec:
  containers:
    - name: my-fancy-model-including-app
      image: myapp:latest-with-model
```

The Problem Before

Container design - either lightweight + limited:

- CloudNativePG

Or bulky:

- StackGres (100+ extensions)

The Problem Before

```
apiVersion: v1
kind: Pod
metadata:
  name: my-favourite-app
spec:
  initContainers:
    # download the dependency here
  containers:
    - name: my-fancy-model-including-app
      image: myapp
      volumeMounts:
        - name: model-file
          mountPath: /mnt/
          readOnly: true
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-favourite-app
spec:
  volumes:
    - name: model-image
      image:
        reference: myregistry/large-base:v1
        pullPolicy: IfNotPresent
  containers:
    - name: my-fancy-model-including-app
      image: my-fancy-app-without-model
      volumeMounts:
        - name: model-image
          mountPath: /mnt/
          readOnly: true
```

Why OCI ImageVolumes Matter

Use Cases:

- **ML Workloads:** Share large model files across containers
- **Microservices:** Common base files and dependencies
- **Addons/Extensions:** Provide as regular image
- **Edge Computing:** Minimize bandwidth usage

Node SWAP Support

Handle memory spikes with flexible memory management

- Status: **Beta** in Kubernetes 1.34 (Burstable QoS only)
- Next: Enhanced controls and monitoring in 1.35
- KEP: [2400](#)

(Off by default)

The Problem Before

No SWAP support.

```
# this fragment goes into the kubelet's configuration file
memorySwap:
  swapBehavior: LimitedSwap
---
# Only non-high-priority Pods under the Burstable QoS tier are permitted to use swap.
apiVersion: v1
kind: Pod
metadata:
  name: my-app
spec:
  containers:
  - name: java-app
    image: openjdk:11
    resources:
      requests:
        memory: "2Gi"
      limits:
        memory: "4Gi"
```

Why SWAP Support Matters

Use Cases:

- **Applications with Dormant Memory:** Java, batch operations, ...
- **Cost optimization**
- **You are in charge** and responsible for your choices

Job Success and Failure Policies

Better job management with fine-grained control

- Status: both **Stable** in Kubernetes 1.34
- KEP: 3329 (failure), 3998 (success)

The Problem Before

```
apiVersion: batch/v1
kind: Job
metadata:
  name: my-favourite-job
spec:
  backoffLimit: 3  # Retry 3 times, that's it
  template:
    spec:
      containers:
        - name: batch-processor
          image: batch-app:v1
      restartPolicy: Never
```

```
apiVersion: batch/v1
kind: Job
metadata:
  name: my-other-job
spec:
  backoffLimit: 5
  template:
    spec:
      containers:
        - name: batch-processor
          image: batch-app:v1
        restartPolicy: Never # required if using .podFailurePolicy
  podFailurePolicy:
    rules:
      - action: FailJob
        onExitCodes:
          containerName: batch-processor
          operator: In
          values: [1, 2] # Fatal errors - don't retry
      - action: Ignore
        onPodConditions:
          - type: DisruptionTarget # Node maintenance - retry
```

```
apiVersion: batch/v1
kind: Job
metadata:
  name: my-matrix-job
spec:
  backoffLimit: 5
  template:
    spec:
      containers:
        - name: batch-processor
          image: batch-app:v1
        restartPolicy: Never
  parallelism: 10
  completions: 10
  completionMode: Indexed
  successPolicy:
    rules:
      - succeededIndexes: "0-9"
        succeededCount: 8 # 80% success rate is enough
```

Why Job Policies Matter

Use Cases:

- **Data Processing:** Skip corrupted data, continue with rest
- **ML Training:** Handle node failures gracefully
- **ETL Pipelines:** Categorize transient vs permanent failures
- **Batch Analytics:** Partial success scenarios

So your jobs don't need to be idempotent.

Dynamic Resource Allocation

Smart allocation of node resources

- Status: **Stable** in Kubernetes 1.34
- Multiple KEPs

```
apiVersion: resource.k8s.io/v1
kind: ResourceSlice
# ...
spec:
  devices:
    - attributes:
        type:
          string: gpu
      capacity:
        memory:
          value: 64Gi
      name: gpu-0
    - attributes:
        type:
          string: gpu
      capacity:
        memory:
          value: 64Gi
      name: gpu-1
  driver: driver.example.com
  nodeName: cluster-1-node-1
```

Why Dynamic Resource Allocation Matters

Use Cases:

- **ML:** Assigning accelerator / graphics cards
- **SR-IOV:** Slicing Network Interface (NICs) cards
- **NVMe Slicing**

What's Next in Kubernetes

Upcoming features to watch:

- **Gang Scheduling** - Coordinate scheduling of related pods ("PodGroup")
 - KEP: [4671](#)
 - Status: Alpha development for 1.35 (?)

Thank You!

*Let's bring more workloads on
Kubernetes!*

Questions?

